

## 📌 What Is the Project?

This is a Python-based simulation of a **lightweight intrusion detection system (IDS)** that uses **entropy** to detect potential **Distributed Denial of Service (DDoS)** attacks based on traffic patterns—specifically, the diversity of IP addresses sending traffic.

It simulates both **normal traffic** (healthy, diverse network activity) and **attack traffic** (overloaded, malicious traffic from a few sources), and applies information theory (entropy calculation) to distinguish between them.

### 1. ✅ First: What is a DDoS Attack?

A **Distributed Denial of Service (DDoS)** attack is when **multiple systems flood a target server or service** with excessive traffic, aiming to overwhelm it and render it unusable.

Typical characteristics:

- Huge volume of requests
- Requests often come from a **small set of IP addresses** (botnets or spoofed)
- Traffic pattern is repetitive and lacks diversity

### 2. 📊 Why Entropy? What Is It?

**Entropy**, in this context, measures the **randomness** or **unpredictability** in a set of values.

From Shannon's Information Theory:

$$H = - \sum p(i) \cdot \log_2 p(i)$$

Where:

- $p(i)$  is the **probability** of each unique IP address in the traffic
- $H$  is the total entropy

More IP diversity = **higher entropy**

Fewer repeated IPs = **lower entropy**

#### 🧠 Intuition:

- A crowd of people from all countries = high entropy (normal traffic)
- 1,000 clones of the same person = low entropy (attack traffic)

### 3. 🛠️ Simulation Logic

This project **simulates traffic** (not actual packet sniffing):

- **Normal traffic:** randomly selected from a wide IP pool → higher entropy
- **Attack traffic:** mostly from a few repeated IPs → lower entropy

You simulate this using Python's `random.choices()` function to generate packets coming from a list of IPs.

## 4. 🔍 The Detection Pipeline

Here's the actual logic pipeline of the code:

Step	Description
<code>send_packets()</code>	Simulates traffic with a uniform (normal) or concentrated (attack) IP distribution
<code>compute_entropy()</code>	Calculates entropy of the traffic based on frequency of each IP
<code>detect_ddos()</code>	Compares entropy to a <b>threshold</b> value to classify the traffic
<code>find_accuracy()</code>	Repeats this process for multiple test cases to measure how accurate the classifier is

## 5. 🎯 What's the Threshold?

You need a **cutoff point** to decide when traffic is suspicious.

The project sets the **threshold entropy** as the average of:

- Normal traffic entropy
- Attack traffic entropy

Then:

- If current traffic entropy < threshold → classify as DDoS
- Otherwise → classify as normal

This approach is **unsupervised**, meaning no labeled training data is needed—it learns the threshold dynamically during simulation.

## 6. ⚙️ Key Parameters You Can Tweak

Parameter	What It Controls
<code>num_packets</code>	Total number of simulated packets in a session
<code>num_ips</code>	Total number of unique IPs to simulate normal diversity
<code>attack=True/False</code>	Whether traffic simulates an attack
<code>num_tests</code>	How many total test iterations are run
<code>threshold_entropy</code>	The cutoff value for detection logic

## 7. 🚧 Accuracy & Evaluation

You run tests with both attack and normal traffic.

Each time:

- It checks whether the classifier labels the traffic correctly.
- Over multiple test cycles, it computes a **detection accuracy score**.

You can also turn on `verbose=True` to **see the classification output** for every test.

## 🌟 Limitations of This Approach

- It's a **simulation** — no real-time packet sniffing
- **Spoofed or distributed attacks** with many IPs could bypass this (entropy evasion)
- The threshold is **hardcoded** from an average—no learning or adaptability
- Doesn't consider **packet rate, payloads, or timestamps**

## 🚩 Phase 1: From Simulation to Real-Time Detection

### Option B: Use Packet Sniffers (If You Want Deep Packet Inspection)

- Tools like `scapy`, `tcpdump`, or `pyshark` can capture traffic directly.
- Risk: requires root access, slows down your server, not ideal for production apps unless absolutely needed.

### 3. Why use entropy instead of rate limiting or signature-based detection?

- **Signature-based:** Fails for novel attacks.
- **Rate limiting:** Works per IP, doesn't catch distributed attacks.
- **Entropy:** Doesn't rely on known patterns, adapts to new threats, works at the aggregate level across IPs.

### 9. Could an attacker bypass this system by spoofing IPs?

Yes. If the attacker uses many spoofed IPs to artificially inflate entropy, the system might classify it as normal. That's why real-world systems combine entropy with behavioral analysis or deep packet inspection.

## ◆ EDGE CASES & STABILITY

### 11. What if traffic is very low (e.g., <100 packets)?

Low sample size can skew entropy. To handle this:

- Require minimum packet count before computing entropy
- Use time-window buffering
- Apply entropy smoothing or normalization

### 12. What happens if normal traffic has low entropy?

That can happen, e.g., a few power users or internal traffic. In that case:

- Adjust the threshold dynamically
- Compare current entropy against historical baseline instead of static threshold

### 13. How does your model handle false positives/negatives?

It doesn't—yet. That's one limitation. To handle them:

- Add confidence intervals
- Introduce a feedback loop for misclassifications
- Use multiple features, not just entropy

### 14. What is the minimum entropy value possible in your system?

If all traffic comes from one IP, entropy = 0.

This represents the most deterministic distribution—complete certainty about the source.

### 15. What kind of real-world data would break this system?

- Load balancers rotating IPs rapidly
  - IoT devices with bursty but uniform traffic
  - CDN nodes all funneling traffic from a few IPs
- These could be flagged as DDoS even if legit.